

[illegible]

6. Implement gradient descent & back propagation in deep neural network

Aim :-

To implement gradient descent and back propagation algorithm in a simple deep neural network and study their role in training.

Objectives

- ① to understand the working of gradient descent as an optimization method
- ② to implement back propagation for updating weights in neural network
- ③ to train single neural network for a classification task using this algorithm

Pseudo Code :-

- ① Start
- ② Initialize dataset (x, y) for training
- ③ Initialize neural network parameters
 - * Input layer, hidden layer, output layer
 - * Random weights & biases
 - * Learning rate

④ Forward propagation

$$z_1 = w_1 * x + b_1$$

A_1 = activation

A_2 = activation z_2

$$z_2 = w_2 * A_1 + b_2$$

⑤ compute loss using mean square error

⑥ Back propagation

* compute error at output

* calculate gradients for w & b

⑦ update weight and bias

$$w = w - \eta * dw$$

$$b = b - \eta * db$$

⑧ Repeat the epoches

⑨ observe loss and accuracy improvement

Observation

→ loss decreases as the number of iteration epoch increases

→ weights and biases adjust to minimize the error

→ Learning rate (η) quantity influences convergence speed.

~~0.001, 0.01, 0.1~~

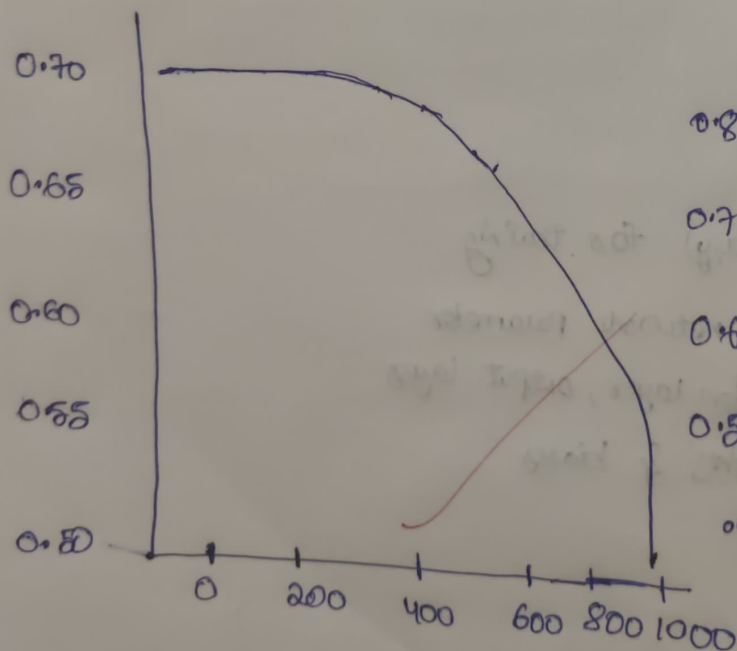
Results

Successfully implemented gradient descent and back propagation in deep learning.

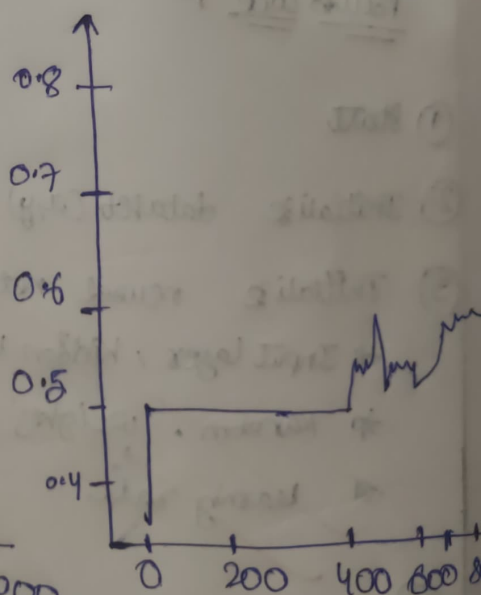
Outputs

epoch 0 ; loss : 0.6932 ; Accuracy : 0.3650
 epoch 200 ; loss : 0.6921 ; Accuracy : 0.5200
 epoch 300 ; loss : 0.6910 ; 0.5200
 epoch 400 ; 0.6849 ; 0.5200
 epoch 500 ; 0.6631 ; 0.5200
 epoch 600 ; 0.6040 ; 0.5300
 epoch 700 ; 0.6250
 epoch 800 ; 0.5729 ; 0.6900
 epoch 900 ; 0.5324 ; 0.8250

Training loss



Training accuracy





```
import numpy as np
import matplotlib.pyplot as plt

# --- Create toy dataset (binary classification) ---
np.random.seed(42)
X = np.random.randn(200, 2)  # 200 samples, 2 features
y = (X[:, 0] * X[:, 1] > 0).astype(int).reshape(-1, 1)

# --- Helper functions ---
def sigmoid(z):
    return 1 / (1 + np.exp(-z))

def relu(z):
    return np.maximum(0, z)

def relu_deriv(z):
    return (z > 0).astype(float)

# --- Initialize weights ---
hidden_units = 4
W1 = np.random.randn(2, hidden_units) * 0.01
b1 = np.zeros((1, hidden_units))
W2 = np.random.randn(hidden_units, 1) * 0.01
b2 = np.zeros((1, 1))
```




```
# --- Initialize weights ---
hidden_units = 4
W1 = np.random.randn(2, hidden_units) * 0.01
b1 = np.zeros((1, hidden_units))
W2 = np.random.randn(hidden_units, 1) * 0.01
b2 = np.zeros((1, 1))

# --- Training loop ---
lr = 0.1
epochs = 1000
losses = []          # store losses
accuracies = []      # store accuracies
epoch_list = []

for epoch in range(epochs):
    # Forward pass
    Z1 = X.dot(W1) + b1
    A1 = relu(Z1)
    Z2 = A1.dot(W2) + b2
    A2 = sigmoid(Z2)

    # Loss (binary cross entropy)
    m = X.shape[0]
    loss = -np.mean(y * np.log(A2 + 1e-8) + (1 - y) * np.log(1 - A2 + 1e-8))
    losses.append(loss)

    # Accuracy
    predictions = (A2 > 0.5).astype(int)
    accuracy = np.mean(predictions == y)
    accuracies.append(accuracy)

    epoch_list.append(epoch)

    # Backpropagation
    dZ2 = A2 - y
    dW2 = (A1.T.dot(dZ2)) / m
    db2 = np.mean(dZ2, axis=0, keepdims=True)

    dA1 = dZ2.dot(W2.T)
    dZ1 = dA1 * relu_deriv(Z1)
    dW1 = (X.T.dot(dZ1)) / m
    db1 = np.mean(dZ1, axis=0, keepdims=True)
```

▶

```
# Gradient descent update
```

```
W1 -= lr * dW1
```

```
b1 -= lr * db1
```

```
W2 -= lr * dW2
```

```
b2 -= lr * db2
```

```
# Print every 100 epochs
```

```
if epoch % 100 == 0:
```

```
    print(f"Epoch {epoch}, Loss: {loss:.4f}, Accuracy: {accuracy:.4f}")
```

```
# --- Plot Loss & Accuracy ---
```

```
plt.figure(figsize=(10,4))
```

```
plt.subplot(1,2,1)
```

```
plt.plot(epoch_list, losses, label="Loss")
```

```
plt.xlabel("Epochs")
```

```
plt.ylabel("Loss")
```

```
plt.title("Training Loss")
```

```
plt.legend()
```

```
plt.subplot(1,2,2)
```

```
plt.plot(epoch_list, accuracies, label="Accuracy", color="orange")
```

```
plt.xlabel("Epochs")
```

```
plt.ylabel("Accuracy")
```

```
plt.title("Training Accuracy")
```

```
plt.legend()
```

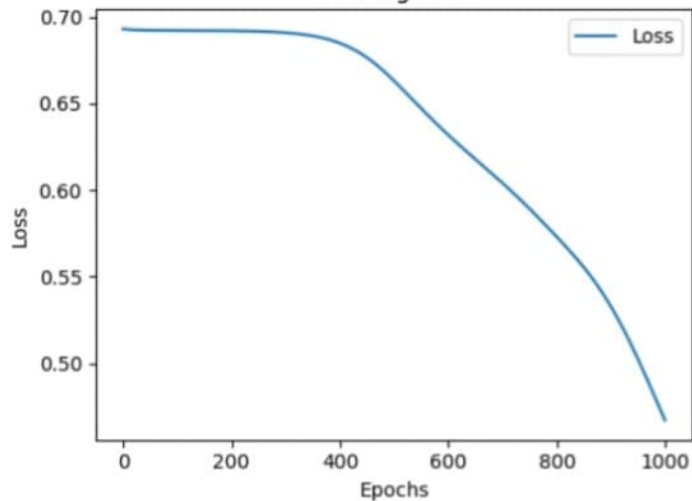
```
plt.tight_layout()
```

```
plt.show()
```



Epoch 0, Loss: 0.6932, Accuracy: 0.3650
Epoch 100, Loss: 0.6923, Accuracy: 0.5200
Epoch 200, Loss: 0.6921, Accuracy: 0.5200
Epoch 300, Loss: 0.6910, Accuracy: 0.5200
Epoch 400, Loss: 0.6849, Accuracy: 0.5200
Epoch 500, Loss: 0.6631, Accuracy: 0.5300
Epoch 600, Loss: 0.6318, Accuracy: 0.6200
Epoch 700, Loss: 0.6040, Accuracy: 0.6350
Epoch 800, Loss: 0.5729, Accuracy: 0.6700
Epoch 900, Loss: 0.5324, Accuracy: 0.8250

Training Loss



Training Accuracy

